

Title: Implement data structures in R on datasets.

Tools: R studio

Theory: A data structure is a particular way of organizing data in a computer so that it can be used effectively. The idea is to reduce the space and time complexities of different tasks. Data structures in R programming are tools for holding multiple values.

R's base data structures are often organized by their dimensionality (1D, 2D, or nD) and whether they're homogeneous (all elements must be of the identical type) or heterogeneous (the elements are often of various types). This gives rise to the six data types which are most frequently utilized in data analysis.

ration with other programming languages like Python, C, and Java.

The most essential data structures used in R include:

- Vectors
- Lists
- Dataframes
- Matrices
- Arrays
- Factors
- Tibbles

- 1) From the mtcars dataset, create a vector from the mpg column. Replace values greater than the 75th percentile with "High" and less than the 25th percentile with "Low", leaving the rest as "Medium".

Code:

```
data(mtcars)
mpg_vector <- mtcars$mpg
lower_percentile <- quantile(mpg_vector, 0.25)
upper_percentile <- quantile(mpg_vector, 0.75)
mpg_category <- ifelse(mpg_vector > upper_percentile, "High",
                      ifelse(mpg_vector < lower_percentile, "Low", "Medium"))
print(mpg_category)
```

Output:

```
data(mtcars)
mpg_vector <- mtcars$mpg
lower_percentile <- quantile(mpg_vector, 0.25)
upper_percentile <- quantile(mpg_vector, 0.75)
mpg_category <- ifelse(mpg_vector > upper_percentile, "High",
                      ifelse(mpg_vector < lower_percentile, "Low", "Medium"))
print(mpg_category)

> data(mtcars)
> mpg_vector <- mtcars$mpg
> lower_percentile <- quantile(mpg_vector, 0.25)
> upper_percentile <- quantile(mpg_vector, 0.75)
> mpg_category <- ifelse(mpg_vector > upper_percentile, "High",
+                       ifelse(mpg_vector < lower_percentile, "Low", "Medium"))
> print(mpg_category)
[1] "Medium" "Medium" "Medium" "Medium" "Medium" "Medium" "Low"   "High"
[9] "Medium" "Medium" "Medium" "Medium" "Medium" "Low"   "Low"   "Low"
[17] "Low"    "High"    "High"    "High"    "Medium" "Medium" "Low"   "Low"
[25] "Medium" "High"    "High"    "High"    "Medium" "Medium" "Low"   "Medium"
> |
```

- 2) Create two vectors from the iris dataset (Sepal.Length and Sepal.Width). Calculate the element-wise difference and product. Find the sum of only the positive differences.

Code:

```
data(iris)
sepal_length <- iris$Sepal.Length
sepal_width <- iris$Sepal.Width
difference <- sepal_length - sepal_width
product <- sepal_length * sepal_width
positive_difference_sum <- sum(difference[difference > 0])
print(paste("Sum of positive differences: ", positive_difference_sum))
print("Product of Sepal.Length and Sepal.Width: ")
print(product)
```

Output:

```
> data(iris)
> sepal_length <- iris$Sepal.Length
> sepal_width <- iris$Sepal.Width
> difference <- sepal_length - sepal_width
> product <- sepal_length * sepal_width
> positive_difference_sum <- sum(difference[difference > 0])
> print(paste("Sum of positive differences: ", positive_difference_sum))
[1] "Sum of positive differences: 417.9"
> print("Product of Sepal.Length and Sepal.Width: ")
[1] "Product of Sepal.Length and Sepal.Width: "
> print(product)
[1] 17.85 14.70 15.04 14.26 18.00 21.06 15.64 17.00 12.76 15.19 19.98 16.32
[13] 14.40 12.90 23.20 25.08 21.06 17.85 21.66 19.38 18.36 18.87 16.56 16.83
[25] 16.32 15.00 17.00 18.20 17.68 15.04 14.88 18.36 21.32 23.10 15.19 16.00
```

```
[37] 19.25 17.64 13.20 17.34 17.50 10.35 14.08 17.50 19.38 14.40 19.38 14.72
[49] 19.61 16.50 22.40 20.48 21.39 12.65 18.20 15.96 20.79 11.76 19.14 14.04
[61] 10.00 17.70 13.20 17.69 16.24 20.77 16.80 15.66 13.64 14.00 18.88 17.08
[73] 15.75 17.08 18.56 19.80 19.04 20.10 17.40 14.82 13.20 13.20 15.66 16.20
[85] 16.20 20.40 20.77 14.49 16.80 13.75 14.30 18.30 15.08 11.50 15.12 17.10
[97] 16.53 17.98 12.75 15.96 20.79 15.66 21.30 18.27 19.50 22.80 12.25 21.17
[109] 16.75 25.92 20.80 17.28 20.40 14.25 16.24 20.48 19.50 29.26 20.02 13.20
[121] 22.08 15.68 21.56 17.01 22.11 23.04 17.36 18.30 17.92 21.60 20.72 30.02
[133] 17.92 17.64 15.86 23.10 21.42 19.84 18.00 21.39 20.77 21.39 15.66 21.76
[145] 22.11 20.10 15.75 19.50 21.08 17.70
```

- 3) From the mtcars dataset, extract rows where mpg is greater than the median and hp is below the median. Find the correlation between mpg and hp for this subset.

Code:

```
data(mtcars)
median_mpg <- median(mtcars$mpg)
median_hp <- median(mtcars$hp)
subset_data <- mtcars[mtcars$mpg > median_mpg & mtcars$hp < median_hp, ]
correlation <- cor(subset_data$mpg, subset_data$hp)
print(correlation)
```

Output:

```
> data(mtcars)
> median_mpg <- median(mtcars$mpg)
> median_hp <- median(mtcars$hp)
> subset_data <- mtcars[mtcars$mpg > median_mpg & mtcars$hp < median_hp, ]
> correlation <- cor(subset_data$mpg, subset_data$hp)
> print(correlation)
[1] -0.6537747
```

- 4) Create a list from the mtcars dataset containing mpg, hp, and cyl. Update the list to include the row names as a new element. Then filter out the cars where mpg is less than 20.

Code:

```
data(mtcars)
car_list <- list(
  mpg = mtcars$mpg,
  hp = mtcars$hp,
  cyl = mtcars$cyl
)
car_list$row_names <- rownames(mtcars)
car_df <- data.frame(car_list)
filtered_cars <- car_df[car_df$mpg >= 20, ]
print(filtered_cars)
```

Output:

```

> data(mtcars)
> car_list <- list(
+   mpg = mtcars$mpg,
+   hp = mtcars$hp,
+   cyl = mtcars$cyl
+ )
> car_list$row_names <- rownames(mtcars)
> car_df <- data.frame(car_list)
> filtered_cars <- car_df[car_df$mpg >= 20, ]
> print(filtered_cars)
   mpg  hp cyl   row_names
1  21.0 110   6   Mazda RX4
2  21.0 110   6 Mazda RX4 Wag
3  22.8  93   4   Datsun 710
4  21.4 110   6 Hornet 4 Drive
8  24.4  62   4    Merc 240D
9  22.8  95   4    Merc 230
18 32.4  66   4    Fiat 128
19 30.4  52   4   Honda Civic
20 33.9  65   4 Toyota Corolla
21 21.5  97   4 Toyota Corona
26 27.3  66   4    Fiat X1-9
27 26.0  91   4  Porsche 914-2
28 30.4 113   4   Lotus Europa
32 21.4 109   4   Volvo 142E

```

- 5) From the iris dataset, create a list containing Sepal.Length, Sepal.Width, and Species. Merge this list with a new list containing a summary of Sepal.Length using the summary() function.

Code:

```

data(iris)
iris_list <- list(
  Sepal.Length = iris$Sepal.Length,
  Sepal.Width = iris$Sepal.Width,
  Species = iris$Species
)
sepal_length_summary <- summary(iris$Sepal.Length)
merged_list <- c(iris_list, Sepal.Length.Summary = sepal_length_summary)
print(merged_list)

```

Output:

```

> data(iris)
> iris_list <- list(
+   Sepal.Length = iris$Sepal.Length,
+   Sepal.Width = iris$Sepal.Width,
+   Species = iris$Species
+ )
> sepal_length_summary <- summary(iris$Sepal.Length)
> merged_list <- c(iris_list, Sepal.Length.Summary = sepal_length_summary)
> print(merged_list)
$Sepal.Length
 [1] 5.1 4.9 4.7 4.6 5.0 5.4 4.6 5.0 4.4 4.9 5.4 4.8 4.8 4.3 5.8 5.7 5.4 5.1
[19] 5.7 5.1 5.4 5.1 4.6 5.1 4.8 5.0 5.0 5.2 5.2 4.7 4.8 5.4 5.2 5.5 4.9 5.0
[37] 5.5 4.9 4.4 5.1 5.0 4.5 4.4 5.0 5.1 4.8 5.1 4.6 5.3 5.0 7.0 6.4 6.9 5.5
[55] 6.5 5.7 6.3 4.9 6.6 5.2 5.0 5.9 6.0 6.1 5.6 6.7 5.6 5.8 6.2 5.6 5.9 6.1

```

- 6) Create a matrix from the mtcars dataset containing the first 10 rows of mpg, hp, and cyl. Add a new row representing a car with mpg = 18, hp = 150, and cyl = 6, and a new column showing the sum of all rows.

Code:

```
data(mtcars)
car_matrix <- as.matrix(mtcars[1:10, c("mpg", "hp", "cyl")])
new_car <- c(18, 150, 6)
car_matrix <- rbind(car_matrix, new_car)
rownames(car_matrix)[nrow(car_matrix)] <- "New_Car"
car_matrix <- cbind(car_matrix, Sum = rowSums(car_matrix))
print(car_matrix)
```

Output:

```
> data(mtcars)
> car_matrix <- as.matrix(mtcars[1:10, c("mpg", "hp", "cyl")])
> new_car <- c(18, 150, 6)
> car_matrix <- rbind(car_matrix, new_car)
> rownames(car_matrix)[nrow(car_matrix)] <- "New_Car"
> car_matrix <- cbind(car_matrix, Sum = rowSums(car_matrix))
> print(car_matrix)
```

	mpg	hp	cyl	Sum
Mazda RX4	21.0	110	6	137.0
Mazda RX4 Wag	21.0	110	6	137.0
Datsun 710	22.8	93	4	119.8
Hornet 4 Drive	21.4	110	6	137.4
Hornet Sportabout	18.7	175	8	201.7
Valiant	18.1	105	6	129.1
Duster 360	14.3	245	8	267.3
New_Car	18	150	6	

- 7) From a matrix created from mtcars (mpg and hp), write a function that accepts row and column indices and returns the value at that position. If the indices are out of bounds, return "Invalid";.

Code:

```
data(mtcars)
car_matrix <- as.matrix(mtcars[, c("mpg", "hp")])
get_value <- function(matrix, row_index, col_index) {
  if (row_index < 1 || row_index > nrow(matrix) || col_index < 1 || col_index >
ncol(matrix)) {
    return("Invalid")
  } else {
    return(matrix[row_index, col_index])
  }
}
```



```

value1 <- get_value(car_matrix, 5, 1)
print(value1)
value2 <- get_value(car_matrix, 10, 2)
print(value2)
invalid_value <- get_value(car_matrix, 15, 1) # Out of bounds
print(invalid_value)

```

Output:

```

> data(mtcars)
> car_matrix <- as.matrix(mtcars[, c("mpg", "hp")])
> get_value <- function(matrix, row_index, col_index) {
+   if (row_index < 1 || row_index > nrow(matrix) || col_index < 1 || col_index >
+     ncol(matrix)) {
+     return("Invalid")
+   } else {
+     return(matrix[row_index, col_index])
+   }
+ }
> value1 <- get_value(car_matrix, 5, 1)
> print(value1)
[1] 18.7
> value2 <- get_value(car_matrix, 10, 2)
~ print(value2)
> print(value2)
[1] 123
> invalid_value <- get_value(car_matrix, 15, 1) # Out of bounds
> print(invalid_value)
[1] 10.4

```

- 8) Create a 3D array from the iris dataset using Sepal.Length, Sepal.Width, and Petal.Length. Find the row-wise mean for each matrix level and store the result in a new matrix.

Code:

```

data(iris)
iris_array <- array(c(iris$Sepal.Length[1:10],
                     iris$Sepal.Width[1:10],
                     iris$Petal.Length[1:10]),
                  dim = c(10, 3, 1))
dimnames(iris_array) <- list(paste("Row", 1:10),
                             c("Sepal.Length", "Sepal.Width", "Petal.Length"),
                             "Level 1")
print("3D Array:")
print(iris_array)
row_means <- apply(iris_array, c(1, 3), mean)
print("Row-wise Means:")
print(row_means)

```

Output:

```
> data(iris)
> iris_array <- array(c(iris$Sepal.Length[1:10],
+                       iris$Sepal.Width[1:10],
+                       iris$Petal.Length[1:10]),
+                     dim = c(10, 3, 1))
> dimnames(iris_array) <- list(paste("Row", 1:10),
+                               c("Sepal.Length", "Sepal.Width", "Petal.Length"),
+                               "Level 1")
> print("3D Array:")
[1] "3D Array:"
> print(iris_array)
, , Level 1

      Sepal.Length Sepal.Width Petal.Length
Row 1           5.1           3.5           1.4
Row 2           4.9           3.0           1.4
Row 3           4.7           3.2           1.3
Row 4           4.6           3.1           1.5
Row 5           5.0           3.6           1.4
Row 6           5.4           3.9           1.7
Row 7           4.6           3.4           1.4
Row 8           5.0           3.4           1.5
Row 9           4.4           2.9           1.4
Row 10          4.9           3.1           1.5

> row_means <- apply(iris_array, c(1, 3), mean)
> print("Row-wise Means:")
[1] "Row-wise Means:"
> print(row_means)
      Level 1
Row 1 3.333333
Row 2 3.100000
Row 3 3.066667
Row 4 3.066667
Row 5 3.333333
Row 6 3.666667
```

- 9) From an array created using mtcars (mpg, hp, cyl), create a function that calculates the mean and median for a specific matrix level. Use apply() for the calculation.

Code:

```
data(mtcars)
mtcars_array <- array(c(mtcars$mpg, mtcars$hp, mtcars$cyl),
                     dim = c(nrow(mtcars), 3, 1))
dimnames(mtcars_array) <- list(rownames(mtcars),
                               c("mpg", "hp", "cyl"),
                               "Level 1")

print("3D Array:")
print(mtcars_array)

calculate_stats <- function(array, level) {
  if (level < 1 || level > dim(array)[3]) {
    return("Invalid level")
  }
}
```

```

    }
    means <- apply(array[, , level], 1, mean)
    medians <- apply(array[, , level], 1, median)
    result <- data.frame(Mean = means, Median = medians)
    return(result)
}

stats_level_1 <- calculate_stats(mtcars_array, 1)
print("Mean and Median for Level 1:")
print(stats_level_1)

```

Output:

```

> data(mtcars)
> mtcars_array <- array(c(mtcars$mpg, mtcars$hp, mtcars$cyl),
+                       dim = c(nrow(mtcars), 3, 1))
> dimnames(mtcars_array) <- list(rownames(mtcars),
+                                c("mpg", "hp", "cyl"),
+                                "Level 1")
> print("3D Array:")
[1] "3D Array:"
> print(mtcars_array)
, , Level 1

```

	mpg	hp	cyl
Mazda RX4	21.0	110	6
Mazda RX4 Wag	21.0	110	6
Datsun 710	22.8	93	4
Hornet 4 Drive	21.4	110	6
Hornet Sportabout	18.7	175	8
Valiant	18.1	105	6
Duster 360	14.3	245	8
Merc 240D	24.4	62	4
Merc 230	22.8	95	4
Merc 280	19.2	123	6
Merc 280C	17.8	123	6
Merc 450SE	16.4	180	8
Merc 450SL	17.3	180	8
Merc 450SLC	15.2	180	8
Cadillac Fleetwood	10.4	205	8
Lincoln Continental	10.4	215	8

```

> calculate_stats <- function(array, level) {
+   if (level < 1 || level > dim(array)[3]) {
+     return("Invalid level")
+   }
+   means <- apply(array[, , level], 1, mean)
+   medians <- apply(array[, , level], 1, median)
+   result <- data.frame(Mean = means, Median = medians)
+   return(result)
+ }
> stats_level_1 <- calculate_stats(mtcars_array, 1)
> print("Mean and Median for Level 1:")
[1] "Mean and Median for Level 1:"
> print(stats_level_1)

```

	Mean	Median
Mazda RX4	45.66667	21.0
Mazda RX4 Wag	45.66667	21.0
Datsun 710	39.93333	22.8
Hornet 4 Drive	45.80000	21.4
Hornet Sportabout	67.23333	18.7
Valiant	43.03333	18.1
Duster 360	89.10000	14.3
Merc 240D	30.13333	24.4
Merc 230	40.60000	22.8
Merc 280	49.40000	19.2
Merc 280C	48.93333	17.8
Merc 450SE	68.13333	16.4
Merc 450SL	68.43333	17.3
Merc 450SLC	67.73333	15.2

- 10) Create a data frame from the iris dataset containing Sepal.Length, Sepal.Width, and Species. Create a function that takes a species name as input and returns the average Sepal.Length for that species.

Code:

```
data(iris)
iris_df <- iris[, c("Sepal.Length", "Sepal.Width", "Species")]
print("Iris Data Frame:")
print(head(iris_df))
average_sepal_length <- function(species_name) {
  filtered_data <- iris_df[iris_df$Species == species_name, ]
  if (nrow(filtered_data) == 0) {
    return("Species not found")
  }
  avg_length <- mean(filtered_data$Sepal.Length)
  return(avg_length)
}
species_to_check <- "setosa"
avg_length_setosa <- average_sepal_length(species_to_check)
print(paste("Average Sepal.Length for", species_to_check, ":",
avg_length_setosa))
species_not_found <- "unknown_species"
avg_length_unknown <- average_
```

Output:

```
> data(iris)
> iris_df <- iris[, c("Sepal.Length", "Sepal.Width", "Species")]
> print("Iris Data Frame:")
[1] "Iris Data Frame:"
> print(head(iris_df))
  Sepal.Length Sepal.Width Species
1          5.1           3.5  setosa
2          4.9           3.0  setosa
3          4.7           3.2  setosa
4          4.6           3.1  setosa
5          5.0           3.6  setosa
6          5.4           3.9  setosa
> average_sepal_length <- function(species_name) {
+   filtered_data <- iris_df[iris_df$Species == species_name, ]
+   if (nrow(filtered_data) == 0) {
+     return("Species not found")
+   }
+   avg_length <- mean(filtered_data$Sepal.Length)
+   return(avg_length)
+ }
> species_to_check <- "setosa"
> avg_length_setosa <- average_sepal_length(species_to_check)
> print(paste("Average Sepal.Length for", species_to_check, ":", avg_length_setos
a))
[1] "Average Sepal.Length for setosa : 5.006"
> species_not_found <- "unknown_species"
> avg_length_unknown <- average_sepal_length(species_not_found)
> print(avg_length_unknown)
[1] "Species not found"
```

11) From the mtcars dataset, create a new column power_to_weight as the ratio of hp to wt. Then find the car with the highest and lowest power_to_weight ratio.

Code:

```
data(mtcars)
mtcars$power_to_weight <- mtcars$hp / mtcars$wt
print("Updated mtcars dataset with power_to_weight:")
print(head(mtcars))
highest_power_to_weight <- mtcars[which.max(mtcars$power_to_weight), ]
lowest_power_to_weight <- mtcars[which.min(mtcars$power_to_weight), ]
print("Car with the highest power_to_weight ratio:")
print(highest_power_to_weight)
print("Car with the lowest power_to_weight ratio:")
print(lowest_power_to_weight)
```

Output:

```
> data(mtcars)
> mtcars$power_to_weight <- mtcars$hp / mtcars$wt
> print("Updated mtcars dataset with power_to_weight:")
[1] "Updated mtcars dataset with power_to_weight:"
> print(head(mtcars))
      mpg  cyl  disp  hp drat   wt  qsec vs  am  gear  carb
Mazda RX4         21.0   6  160 110 3.90 2.620 16.46 0   1    4    4
Mazda RX4 Wag     21.0   6  160 110 3.90 2.875 17.02 0   1    4    4
Datsun 710         22.8   4  108  93 3.85 2.320 18.61 1   1    4    1
Hornet 4 Drive     21.4   6  258 110 3.08 3.215 19.44 1   0    3    1
Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02 0   0    3    2
Valiant           18.1   6  225 105 2.76 3.460 20.22 1   0    3    1
      power_to_weight
Mazda RX4           41.98473
Mazda RX4 Wag       38.26087
Datsun 710           40.08621
Hornet 4 Drive       34.21462
Hornet Sportabout    50.87209
Valiant              30.34682
> highest_power_to_weight <- mtcars[which.max(mtcars$power_to_weight), ]
> lowest_power_to_weight <- mtcars[which.min(mtcars$power_to_weight), ]
> print("Car with the highest power_to_weight ratio:")
[1] "Car with the highest power_to_weight ratio:"
> print(highest_power_to_weight)
      mpg  cyl  disp  hp drat   wt  qsec vs  am  gear  carb power_to_weight
Maserati Bora  15   8  301 335 3.54 3.57 14.6 0   1    5    8      93.83754
> print("Car with the lowest power_to_weight ratio:")
[1] "Car with the lowest power_to_weight ratio:"
> print(lowest_power_to_weight)
      mpg  cyl  disp  hp drat   wt  qsec vs  am  gear  carb power_to_weight
Merc 240D  24.4   4 146.7  62 3.69 3.19  20  1  0    4    2      19.43574
```

12) From the iris dataset, create a new data frame where the Sepal.Length is greater than 5.5. Then calculate the correlation matrix for the numeric columns.

Code:

```
data(iris)
iris_filtered <- iris[iris$Sepal.Length > 5.5, ]
print("Filtered Iris Data Frame (Sepal.Length > 5.5):")
print(head(iris_filtered))
correlation_matrix <- cor(iris_filtered[, sapply(iris_filtered, is.numeric)])
print("Correlation Matrix for Numeric Columns:")
print(correlation_matrix)
```

Output:

```
> data(iris)
> iris_filtered <- iris[iris$Sepal.Length > 5.5, ]
> print("Filtered Iris Data Frame (Sepal.Length > 5.5):")
[1] "Filtered Iris Data Frame (Sepal.Length > 5.5):"
> print(head(iris_filtered))
   Sepal.Length Sepal.Width Petal.Length Petal.Width   Species
15          5.8         4.0         1.2         0.2    setosa
16          5.7         4.4         1.5         0.4    setosa
19          5.7         3.8         1.7         0.3    setosa
51          7.0         3.2         4.7         1.4 versicolor
52          6.4         3.2         4.5         1.5 versicolor
53          6.9         3.1         4.9         1.5 versicolor
> correlation_matrix <- cor(iris_filtered[, sapply(iris_filtered, is.numeric)])
> print("Correlation Matrix for Numeric Columns:")
[1] "Correlation Matrix for Numeric Columns:"
> print(correlation_matrix)
               Sepal.Length Sepal.Width Petal.Length Petal.Width
Sepal.Length   1.0000000   0.21495443   0.7121662   0.52271575
Sepal.Width    0.2149544   1.00000000  -0.1377127   0.01925592
Petal.Length   0.7121662  -0.13771270   1.0000000   0.85330666
Petal.Width    0.5227157   0.01925592   0.8533067   1.00000000
```

14) Create a factor from the Species column of iris. Create a contingency table showing the frequency of each species combined with the Sepal.Width above and below the median.

Code:

```
data(iris)
iris$Species <- as.factor(iris$Species)
median_sepal_width <- median(iris$Sepal.Width)
iris$Width_Category <- ifelse(iris$Sepal.Width > median_sepal_width, "Above
Median", "Below Median")
contingency_table <- table(iris$Species, iris$Width_Category)
print("Contingency Table of Species and Sepal.Width Category:")
print(contingency_table)
```

Output:


```

> data(iris)
> iris$Species <- as.factor(iris$Species)
> median_sepal_width <- median(iris$Sepal.Width)
> iris$width_Category <- ifelse(iris$Sepal.Width > median_sepal_width, "Above Median", "Below Median")
> contingency_table <- table(iris$Species, iris$width_Category)
> print("Contingency Table of Species and Sepal.Width Category:")
[1] "Contingency Table of Species and Sepal.Width Category:"
> print(contingency_table)

```

	Above Median	Below Median
setosa	42	8
versicolor	8	42
virginica	17	33

15) From the mtcars dataset, identify rows with missing values. Replace missing values in mpg and hp with the median of their respective columns.

Code:

```

data(mtcars)
missing_values <- is.na(mtcars)
rows_with_missing <- mtcars[apply(missing_values, 1, any), ]
print("Rows with missing values:")
print(rows_with_missing)
median_mpg <- median(mtcars$mpg, na.rm = TRUE)
median_hp <- median(mtcars$hp, na.rm = TRUE)
mtcars$mpg[is.na(mtcars$mpg)] <- median_mpg
mtcars$hp[is.na(mtcars$hp)] <- median_hp
print("Updated mtcars dataset after replacing missing values:")
print(head(mtcars))

```

Output:

```

> data(mtcars)
> missing_values <- is.na(mtcars)
> rows_with_missing <- mtcars[apply(missing_values, 1, any), ]
> print("Rows with missing values:")
[1] "Rows with missing values:"
> print(rows_with_missing)
[1] mpg cyl disp hp drat wt  qsec vs am gear carb
<0 rows> (or 0-length row.names)
> median_mpg <- median(mtcars$mpg, na.rm = TRUE)
> median_hp <- median(mtcars$hp, na.rm = TRUE)
> mtcars$mpg[is.na(mtcars$mpg)] <- median_mpg
> mtcars$hp[is.na(mtcars$hp)] <- median_hp
> print("Updated mtcars dataset after replacing missing values:")
[1] "Updated mtcars dataset after replacing missing values:"
> print(head(mtcars))

```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

17) Create a function that takes a data frame and a column name as input, and converts continuous numeric data into three categories: "Low", "Medium", "High", based on quantiles. Test it on mtcars\$hp.

Code:

```

categorize_continuous <- function(data, column_name) {
  if (!column_name %in% names(data)) {
    stop("Column not found in the data frame.")
  }
  column_data <- data[[column_name]]
  quantiles <- quantile(column_data, probs = c(0, 1/3, 2/3, 1), na.rm = TRUE)
  categories <- cut(column_data,
    breaks = quantiles,
    labels = c("Low", "Medium", "High"),
    include.lowest = TRUE)

  return(categories)
}

mtcars$hp_category <- categorize_continuous(mtcars, "hp")
print("Updated mtcars dataset with hp categories:")
print(head(mtcars))

```

Output:

```

> categorize_continuous <- function(data, column_name) {
+   if (!column_name %in% names(data)) {
+     stop("Column not found in the data frame.")
+   }
+   column_data <- data[[column_name]]
+   quantiles <- quantile(column_data, probs = c(0, 1/3, 2/3, 1), na.rm = TRUE)
+   categories <- cut(column_data,
+     breaks = quantiles,
+     labels = c("Low", "Medium", "High"),
+     include.lowest = TRUE)
+   return(categories)
+ }
> mtcars$hp_category <- categorize_continuous(mtcars, "hp")
> print("Updated mtcars dataset with hp categories:")
[1] "Updated mtcars dataset with hp categories:"
> print(head(mtcars))
  mpg  cyl  disp  hp  drat    wt  qsec vs  am  gear  carb
Mazda RX4         21.0   6  160  110  3.90  2.620 16.46  0   1    4    4
Mazda RX4 Wag     21.0   6  160  110  3.90  2.875 17.02  0   1    4    4
Datsun 710        22.8   4  108   93  3.85  2.320 18.61  1   1    4    1
Hornet 4 Drive    21.4   6  258  110  3.08  3.215 19.44  1   0    3    1
Hornet Sportabout 18.7   8  360  175  3.15  3.440 17.02  0   0    3    2
Valiant          18.1   6  225  105  2.76  3.460 20.22  1   0    3    1
  hp_category
Mazda RX4         Medium
Mazda RX4 Wag     Medium
Datsun 710        Low
Hornet 4 Drive    Medium
Hornet Sportabout Medium
Valiant          Low

```

18) From the mtcars dataset, calculate the mean and median mpg for each cyl group using dplyr. Then plot a bar chart showing the differences.

Code:

```

library(dplyr)
library(ggplot2)
data(mtcars)
mpg_summary <- mtcars %>%
  group_by(cyl) %>%

```



```

summarise(
  Mean_MPG = mean(mpg, na.rm = TRUE),
  Median_MPG = median(mpg, na.rm = TRUE)
)
print("Mean and Median MPG for each cyl group:")
print(mpg_summary)

```

Output:

```

> library(dplyr)

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':
  filter, lag

The following objects are masked from 'package:base':
  intersect, setdiff, setequal, union

Warning message:
package 'dplyr' was built under R version 4.4.3
> library(ggplot2)
Warning message:
package 'ggplot2' was built under R version 4.4.3
> data(mtcars)
> mpg_summary <- mtcars %>%
+   group_by(cyl) %>%
+   summarise(
+     Mean_MPG = mean(mpg, na.rm = TRUE),
+     Median_MPG = median(mpg, na.rm = TRUE)
+   )
> print("Mean and Median MPG for each cyl group:")
[1] "Mean and Median MPG for each cyl group:"
> print(mpg_summary)
# A tibble: 3 x 3
   cyl Mean_MPG Median_MPG
  <dbl>   <dbl>     <dbl>
1     4    26.7       26
2     6    19.7      19.7
3     8    15.1      15.2

```

19) Create a function that takes two numeric vectors and returns a list containing the sum, product, difference, and element-wise maximum. Test it on `mtcars$mpg` and `mtcars$hp`.

```

Code: calculate_vector_operations <- function(vec1, vec2) {
  if (length(vec1) != length(vec2)) {
    stop("Vectors must be of the same length.")
  }
  result <- list(
    Sum = vec1 + vec2,
    Product = vec1 * vec2,
    Difference = vec1 - vec2,
    Elementwise_Max = pmax(vec1, vec2)
  )
}

```

```

return(result)
}
result <- calculate_vector_operations(mtcars$mpg, mtcars$hp)
print("Results of vector operations on mtcars$mpg and mtcars$hp:")
print(result)

```

Output:

```

> calculate_vector_operations <- function(vec1, vec2) {
+   if (length(vec1) != length(vec2)) {
+     stop("Vectors must be of the same length.")
+   }
+   result <- list(
+     Sum = vec1 + vec2,
+     Product = vec1 * vec2,
+     Difference = vec1 - vec2,
+     Elementwise_Max = pmax(vec1, vec2)
+   )
+   return(result)
+ }
> result <- calculate_vector_operations(mtcars$mpg, mtcars$hp)
> print("Results of vector operations on mtcars$mpg and mtcars$hp:")
[1] "Results of vector operations on mtcars$mpg and mtcars$hp:"
> print(result)
$Sum
 [1] 131.0 131.0 115.8 131.4 193.7 123.1 259.3  86.4 117.8 142.2 140.8 196.4
[13] 197.3 195.2 215.4 225.4 244.7  98.4  82.4  98.9 118.5 165.5 165.2 258.3
[25] 194.2  93.3 117.0 143.4 279.8 194.7 350.0 130.4

$Product
 [1] 2310.0 2310.0 2120.4 2354.0 3272.5 1900.5 3503.5 1512.8 2166.0 2361.6
[11] 2189.4 2952.0 3114.0 2736.0 2132.0 2236.0 3381.0 2138.4 1580.8 2203.5
[21] 2085.5 2325.0 2280.0 3258.5 3360.0 1801.8 2366.0 3435.2 4171.2 3447.5
[31] 5025.0 2332.6

$Difference
 [1] -89.0 -89.0 -70.2 -88.6 -156.3 -86.9 -230.7 -37.6 -72.2 -103.8
[11] -105.2 -163.6 -162.7 -164.8 -194.6 -204.6 -215.3 -33.6 -21.6 -31.1
[21] -75.5 -134.5 -134.8 -231.7 -155.8 -38.7 -65.0 -82.6 -248.2 -155.3
[31] -320.0 -87.6

$Elementwise_Max
 [1] 110 110  93 110 175 105 245  62  95 123 123 180 180 180 205 215 230  66
[19]  52  65  97 150 150 245 175  66  91 113 264 175 335 109

```

20) Write a function that takes a matrix and returns the row with the highest sum of values. Test it on a matrix created from mtcars\$mpg, mtcars\$hp, and mtcars\$cyl.

Code:

```

calculate_vector_operations <- function(vec1, vec2) {
  if (length(vec1) != length(vec2)) {
    stop("Vectors must be of the same length.")
  }
  result <- list(

```

```

Sum = vec1 + vec2,
Product = vec1 * vec2,
Difference = vec1 - vec2,
Elementwise_Max = pmax(vec1, vec2)
)

return(result)
}
result <- calculate_vector_operations(mtcars$mpg, mtcars$hp)
print("Results of vector operations on mtcars$mpg and mtcars$hp:")
print(result)
row_with_highest_sum <- function(mat) {
  row_sums <- rowSums(mat)
  max_row_index <- which.max(row_sums)
  return(mat[max_row_index, , drop = FALSE]) # drop = FALSE to keep it as a matrix
}
mtcars_matrix <- as.matrix(mtcars[, c("mpg", "hp", "cyl")])
highest_sum_row <- row_with_highest_sum(mtcars_matrix)
print("Row with the highest sum of values:")
print(highest_sum_row)

```

Output:

```

> calculate_vector_operations <- function(vec1, vec2) {
+   if (length(vec1) != length(vec2)) {
+     stop("Vectors must be of the same length.")
+   }
+   result <- list(
+     Sum = vec1 + vec2,
+     Product = vec1 * vec2,
+     Difference = vec1 - vec2,
+     Elementwise_Max = pmax(vec1, vec2)
+   )
+   return(result)
+ }
> result <- calculate_vector_operations(mtcars$mpg, mtcars$hp)
> print("Results of vector operations on mtcars$mpg and mtcars$hp:")
[1] "Results of vector operations on mtcars$mpg and mtcars$hp:"
> print(result)
$Sum
 [1] 131.0 131.0 115.8 131.4 193.7 123.1 259.3  86.4 117.8 142.2 140.8 196.4
[13] 197.3 195.2 215.4 225.4 244.7  98.4  82.4  98.9 118.5 165.5 165.2 258.3
[25] 194.2  93.3 117.0 143.4 279.8 194.7 350.0 130.4

$Product
 [1] 2310.0 2310.0 2120.4 2354.0 3272.5 1900.5 3503.5 1512.8 2166.0 2361.6
[11] 2189.4 2952.0 3114.0 2736.0 2132.0 2236.0 3381.0 2138.4 1580.8 2203.5
[21] 2085.5 2325.0 2280.0 3258.5 3360.0 1801.8 2366.0 3435.2 4171.2 3447.5
[31] 5025.0 2332.6

$Difference
 [1] -89.0 -89.0 -70.2 -88.6 -156.3 -86.9 -230.7 -37.6 -72.2 -103.8
[11] -105.2 -163.6 -162.7 -164.8 -194.6 -204.6 -215.3 -33.6 -21.6 -31.1
[21] -75.5 -134.5 -134.8 -231.7 -155.8 -38.7 -65.0 -82.6 -248.2 -155.3
[31] -320.0 -87.6

$Elementwise_Max
 [1] 110 110  93 110 175 105 245  62  95 123 123 180 180 180 205 215 230  66

```

21) Create a data frame from mtcars and iris. Use merge() to combine them based on a common column name. Handle cases where the column names do not match.

Code:

```

library(dplyr)
data(mtcars)
data(iris)

```

```
mtcars$common_col <- substr(rownames(mtcars), 1, 1) # First letter of car names
iris$common_col <- substr(iris$Species, 1, 1) # First letter of species names
print("mtcars dataset:")
print(head(mtcars))
print("iris dataset:")
print(head(iris))
merged_data <- merge(mtcars, iris, by = "common_col", all = TRUE)
print("Merged dataset:")
```

Output:

```
> library(dplyr)
> data(mtcars)
> data(iris)
> mtcars$common_col <- substr(rownames(mtcars), 1, 1) # First letter of car name
s
> iris$common_col <- substr(iris$Species, 1, 1) # First letter of species names
> print("mtcars dataset:")
[1] "mtcars dataset:"
> print(head(mtcars))
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

```
common_col
Mazda RX4          M
Mazda RX4 Wag      M
Datsun 710         D
Hornet 4 Drive     H
Hornet Sportabout  H
Valiant            V
> print("iris dataset:")
[1] "iris dataset:"
> print(head(iris))
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species	common_col
1	5.1	3.5	1.4	0.2	setosa	s
2	4.9	3.0	1.4	0.2	setosa	s
3	4.7	3.2	1.3	0.2	setosa	s
4	4.6	3.1	1.5	0.2	setosa	s
5	5.0	3.6	1.4	0.2	setosa	s
6	5.4	3.9	1.7	0.4	setosa	s

```
> merged_data <- merge(mtcars, iris, by = "common_col", all = TRUE)
> print("Merged dataset:")
[1] "Merged dataset:"
```

22) Create a data frame from the mtcars dataset. Find the car with the highest and lowest mpg for each cyl group using dplyr and store the result in a new data frame.

Code:

```
library(dplyr)
data(mtcars)
mpg_summary <- mtcars %>%
  group_by(cyl) %>%
  summarise(
    highest_mpg = max(mpg),
```



```

lowest_mpg = min(mpg),
highest_mpg_car = rownames(mtcars[which.max(mpg), ]),
lowest_mpg_car = rownames(mtcars[which.min(mpg), ])
)
print("Summary of highest and lowest mpg for each cyl group:")
print(mpg_summary)

```

Output:

```

> library(dplyr)
> data(mtcars)
> mpg_summary <- mtcars %>%
+   group_by(cyl) %>%
+   summarise(
+     highest_mpg = max(mpg),
+     lowest_mpg = min(mpg),
+     highest_mpg_car = rownames(mtcars[which.max(mpg), ]),
+     lowest_mpg_car = rownames(mtcars[which.min(mpg), ])
+   )
> print("Summary of highest and lowest mpg for each cyl group:")
[1] "Summary of highest and lowest mpg for each cyl group:"
> print(mpg_summary)
# A tibble: 3 × 5
  cyl highest_mpg lowest_mpg highest_mpg_car lowest_mpg_car
  <dbl>      <dbl>      <dbl> <chr>          <chr>
1     4       33.9       21.4 Valiant        Merc 280C
2     6       21.4       17.8 Datsun 710     Valiant
3     8       19.2       10.4 Merc 450SE     Valiant

```

23) Create a matrix from the iris dataset. Write a function that accepts a matrix and returns the row and column indices of the minimum and maximum values.

Code:

```

data(iris)
iris_matrix <- as.matrix(iris[, -5]) # Exclude the Species column
find_min_max_indices <- function(mat) {
  min_value <- min(mat)
  min_indices <- which(mat == min_value, arr.ind = TRUE)
  max_value <- max(mat)
  max_indices <- which(mat == max_value, arr.ind = TRUE)
  return(list(
    min_value = min_value,
    min_indices = min_indices,
    max_value = max_value,
    max_indices = max_indices
  ))
}

```



```
result <- find_min_max_indices(iris_matrix)
print("Minimum value and its indices:")
print(result$min_value)
print(result$min_indices)
print("Maximum value and its indices:")
print(result$max_value)
print(result$max_indices)
```

Output:

```
> data(iris)
> iris_matrix <- as.matrix(iris[, -5]) # Exclude the Species column
> find_min_max_indices <- function(mat) {
+   min_value <- min(mat)
+   min_indices <- which(mat == min_value, arr.ind = TRUE)
+   max_value <- max(mat)
+   max_indices <- which(mat == max_value, arr.ind = TRUE)
+   return(list(
+     min_value = min_value,
+     min_indices = min_indices,
+     max_value = max_value,
+     max_indices = max_indices
+   ))
+ }
> result <- find_min_max_indices(iris_matrix)
> print("Minimum value and its indices:")
[1] "Minimum value and its indices:"
> print(result$min_value)
[1] 0.1
> print(result$min_indices)
      row col
[1,]  10   4
[2,]  13   4
[3,]  14   4
[4,]  33   4
[5,]  38   4
> print("Maximum value and its indices:")
[1] "Maximum value and its indices:"
> print(result$max_value)
[1] 7.9
> print(result$max_indices)
      row col
[1,] 132   1
```

24) Create a correlation matrix for mtcars using cor(). Identify the two columns with the highest and lowest correlation and display the result in a structured format.

Code:

```
data(mtcars)
correlation_matrix <- cor(mtcars)
print("Correlation Matrix:")
print(correlation_matrix)
upper_tri <- correlation_matrix[upper.tri(correlation_matrix)]
highest_correlation <- max(upper_tri)
```

```

lowest_correlation <- min(upper_tri)
highest_indices <- which(correlation_matrix == highest_correlation, arr.ind = TRUE)
lowest_indices <- which(correlation_matrix == lowest_correlation, arr.ind = TRUE)
cat("Highest Correlation:\n")
cat("Columns:", colnames(mtcars)[highest_indices[1, 1]], "and",
    colnames(mtcars)[highest_indices[1, 2]], "\n")
cat("Correlation Value:", highest_correlation, "\n\n")
cat("Lowest Correlation:\n")
cat("Columns:", colnames(mtcars)[lowest_indices[1, 1]], "and",
    colnames(mtcars)[lowest_indices[1, 2]], "\n")
cat("Correlation Value:", lowest_correlation, "\n")

```

Output:

```

> data(iris)
> iris_matrix <- as.matrix(iris[, -5]) # Exclude the Species column
> find_min_max_indices <- function(mat) {
+   min_value <- min(mat)
+   min_indices <- which(mat == min_value, arr.ind = TRUE)
+   max_value <- max(mat)
+   max_indices <- which(mat == max_value, arr.ind = TRUE)
+   return(list(
+     min_value = min_value,
+     min_indices = min_indices,
+     max_value = max_value,
+     max_indices = max_indices
+   ))
+ }
> result <- find_min_max_indices(iris_matrix)
> print("Minimum value and its indices:")
> print(result$min_value)
[1] 0.1
> print(result$min_indices)
  row col
[1,] 10  4
[2,] 13  4
[3,] 14  4
[4,] 33  4
[5,] 38  4
> print("Maximum value and its indices:")
> print(result$max_value)
[1] 7.9
> print(result$max_indices)
  row col
[1,] 132  1
> data(mtcars)
> correlation_matrix <- cor(mtcars)
> print("Correlation Matrix:")
> print(correlation_matrix)
      mpg      cyl      disp      hp      drat      wt      qsec      vs      am      carb
mpg  1.0000000 -0.8521620 -0.8475514 -0.7761684 -0.68117191 -0.8676594
cyl  -0.8521620  1.0000000 -0.9020329 -0.8324475 -0.69993811  0.7824958
disp  -0.8475514  -0.9020329  1.0000000  0.7909486 -0.71021393  0.6879799
hp    -0.7761684  0.8324475  0.7909486  1.0000000 -0.44875912  0.6587479
drat  -0.68117191 -0.69993811 -0.7102139 -0.4487591  1.00000000 -0.7124406
wt    -0.8676594  0.7824958  0.6879799  0.6587479 -0.71244065  1.0000000
qsec  -0.4186840  -0.5912421 -0.4336979 -0.7082234  0.09120476 -0.1747159
vs    -0.6640389  -0.8108118 -0.7104159 -0.7230967  0.44027846 -0.5549157
am    -0.588324  -0.5226070 -0.5912270 -0.2432043  0.71271113 -0.6924953
carb  -0.4802848 -0.4926866 -0.555692  -0.1257043  0.69961013 -0.5832870
      qsec      vs      am      gear      carb
mpg  0.41868403  0.6640389  0.59983243  0.4802848  -0.55092507
cyl  -0.59124207 -0.8108118 -0.52260705 -0.4926866  0.52698829
disp -0.43369788 -0.7104159 -0.59122704 -0.555692  0.39497686
hp    0.70822339 -0.7230967 -0.24320426 -0.1257043  0.74981247
drat  0.09120476  0.4402785  0.71271113  0.6996101 -0.09078980
wt    -0.17471588 -0.5549157 -0.69249526 -0.5832870  0.42760594
qsec  1.00000000  0.7445354 -0.22986086 -0.2126822 -0.65624923
vs    0.74453544  1.0000000  0.16834512  0.2060233 -0.56960714
am    -0.22986086  0.1683451  1.00000000  0.7940588  0.05753435
gear  -0.21268223  0.2060233  0.79405876  1.0000000  0.27407284
carb  -0.65624923 -0.5696071  0.05753435  0.2740728  1.00000000
> upper_tri <- correlation_matrix[upper.tri(correlation_matrix)]
> highest_correlation <- max(upper_tri)
> lowest_correlation <- min(upper_tri)
> highest_indices <- which(correlation_matrix == highest_correlation, arr.ind = TRUE)
> lowest_indices <- which(correlation_matrix == lowest_correlation, arr.ind = TRUE)
> cat("Highest Correlation:\n")
Highest Correlation:
> cat("Columns:", colnames(mtcars)[highest_indices[1, 1]], "and", colnames(mtcars)
[1, 2]), "\n")
Columns: disp and cyl
> cat("Correlation Value:", highest_correlation, "\n\n")
Correlation Value: 0.9020329

> cat("Lowest Correlation:\n")
Lowest Correlation:
> cat("Columns:", colnames(mtcars)[lowest_indices[1, 1]], "and", colnames(mtcars)
[lowest_indices[1, 2]], "\n")
Columns: wt and mpg
> cat("Correlation Value:", lowest_correlation, "\n")
Correlation Value: -0.8676594

```

25) From the iris dataset, create a bar chart showing the average Sepal.Length for each species using ggplot2. Add error bars showing the standard deviation.

Code:

```
library(ggplot2)
library(dplyr)
data(iris)
summary_data <- iris %>%
  group_by(Species) %>%
  summarise(
    avg_sepal_length = mean(Sepal.Length),
    sd_sepal_length = sd(Sepal.Length)
  )
ggplot(summary_data, aes(x = Species, y = avg_sepal_length)) +
  geom_bar(stat = "identity", fill = "skyblue", width = 0.7) +
  geom_errorbar(aes(ymin = avg_sepal_length - sd_sepal_length,
                    ymax = avg_sepal_length + sd_sepal_length),
                width = 0.2) +
  labs(title = "Average Sepal Length by Species",
        x = "Species",
        y = "Average Sepal Length") +
  theme_minimal()
```

Output:

```
> library(ggplot2)
> library(dplyr)
> data(iris)
> summary_data <- iris %>%
+   group_by(Species) %>%
+   summarise(
+     avg_sepal_length = mean(Sepal.Length),
+     sd_sepal_length = sd(Sepal.Length)
+   )
> ggplot(summary_data, aes(x = Species, y = avg_sepal_length)) +
+   geom_bar(stat = "identity", fill = "skyblue", width = 0.7) +
+   geom_errorbar(aes(ymin = avg_sepal_length - sd_sepal_length,
+                     ymax = avg_sepal_length + sd_sepal_length),
+                 width = 0.2) +
+   labs(title = "Average Sepal Length by Species",
+         x = "Species",
+         y = "Average Sepal Length") +
+   theme_minimal()
```



26) Create a matrix from the mtcars dataset. Write a function that accepts a matrix and a row index, and returns the sum of that row. Use apply() for implementation.

Code:

```
data(mtcars)
mtcars_matrix <- as.matrix(mtcars)
row_sum <- function(mat, row_index) {
  if (row_index < 1 || row_index > nrow(mat)) {
    stop("Invalid row index.")
  }
  row_sum_value <- apply(mat[row_index, , drop = FALSE], 1, sum)

  return(row_sum_value)
}
row_index <- 1 # Specify the row index you want to sum
result <- row_sum(mtcars_matrix, row_index)
cat("Sum of row", row_index, "in the mtcars matrix:", result, "\n")
```

Output:


```

> library(ggplot2)
> library(dplyr)
> data(iris)
> summary_data <- iris %>%
+   group_by(Species) %>%
+   summarise(
+     avg_sepal_length = mean(Sepal.Length),
+     sd_sepal_length = sd(Sepal.Length)
+   )
> ggplot(summary_data, aes(x = Species, y = avg_sepal_length)) +
+   geom_bar(stat = "identity", fill = "skyblue", width = 0.7) +
+   geom_errorbar(aes(ymin = avg_sepal_length - sd_sepal_length,
+                     ymax = avg_sepal_length + sd_sepal_length),
+                 width = 0.2) +
+   labs(title = "Average Sepal Length by Species",
+         x = "Species",
+         y = "Average Sepal Length") +
+   theme_minimal()
> data(mtcars)
> mtcars_matrix <- as.matrix(mtcars)
> row_sum <- function(mat, row_index) {
+   if (row_index < 1 || row_index > nrow(mat)) {
+     stop("Invalid row index.")
+   }
+   row_sum_value <- apply(mat[row_index, , drop = FALSE], 1, sum)
+   return(row_sum_value)
+ }
> row_index <- 1 # Specify the row index you want to sum
> result <- row_sum(mtcars_matrix, row_index)
> cat("Sum of row", row_index, "in the mtcars matrix:", result, "\n")
Sum of row 1 in the mtcars matrix: 328.98

```

27) Create a 3D array using iris\$Sepal.Length, iris\$Sepal.Width, and iris\$Petal.Length.

Write a function that returns the mean for each layer and each row in the array.

Code:

```

data(iris)
iris_array <- array(c(iris$Sepal.Length, iris$Sepal.Width, iris$Petal.Length),
                    dim = c(3, 50, 3), # 3 rows (for Sepal.Length, Sepal.Width, Petal.Length),
                    50 samples, 3 species
                    dimnames = list(c("Sepal.Length", "Sepal.Width", "Petal.Length"),
                                     NULL,
                                     levels(iris$Species)))
mean_layer_row <- function(arr) {
  layer_means <- apply(arr, 3, mean) # Mean across the third dimension (layers)
}

```



```
row_means <- apply(arr, 1, mean) # Mean across the first dimension (rows)
```

```
  return(list(layer_means = layer_means, row_means = row_means))
}
```

```
result <- mean_layer_row(iris_array)
```

```
cat("Mean for each layer (species):\n")
```

```
print(result$layer_means)
```

```
cat("\nMean for each row (measurements):\n")
```

```
print(result$row_means)
```

Output:

```
> data(iris)
> iris_array <- array(c(iris$Sepal.Length, iris$Sepal.Width, iris$Petal.Length),
+                      dim = c(3, 50, 3), # 3 rows (for Sepal.Length, Sepal.Width,
+                      Petal.Length), 50 samples, 3 species
+                      dimnames = list(c("Sepal.Length", "Sepal.Width", "Petal.Length"),
+                      NULL,
+                      levels(iris$Species)))
> mean_layer_row <- function(arr) {
+   layer_means <- apply(arr, 3, mean) # Mean across the third dimension (layers)
+   row_means <- apply(arr, 1, mean) # Mean across the first dimension (rows)
+   return(list(layer_means = layer_means, row_means = row_means))
+ }
> result <- mean_layer_row(iris_array)
> cat("Mean for each layer (species):\n")
Mean for each layer (species):
> print(result$layer_means)
   setosa versicolor virginica
5.843333  3.057333  3.758000
> cat("\nMean for each row (measurements):\n")

Mean for each row (measurements):
> print(result$row_means)
Sepal.Length Sepal.Width Petal.Length
   4.200667    4.228000    4.230000
```

28) From the mtcars dataset, create a scatter plot showing the relationship between mpg and hp using ggplot2. Add a regression line with confidence intervals.

Code:

```
library(ggplot2)
```

```
data(mtcars)
```

```
ggplot(mtcars, aes(x = hp, y = mpg)) +
```

```
  geom_point(color = "blue", size = 3) +
```

```
  geom_smooth(method = "lm", se = TRUE, color = "red") +
```

```
  labs(title = "Scatter Plot of MPG vs Horsepower",
```


```
        x = "Horsepower (hp)",
```

```
        y = "Miles per Gallon (mpg)") +
```

```
theme_minimal()
```

Output:

```
> library(ggplot2)
> data(mtcars)
> ggplot(mtcars, aes(x = hp, y = mpg)) +
+   geom_point(color = "blue", size = 3) +
+   geom_smooth(method = "lm", se = TRUE, color = "red") +
+   labs(title = "Scatter Plot of MPG vs Horsepower",
+         x = "Horsepower (hp)",
+         y = "Miles per Gallon (mpg)") +
+   theme_minimal()
`geom_smooth()` using formula = 'y ~ x'
```



29) Create a summary report for the mtcars dataset using rmarkdown. Include a table of descriptive statistics and a correlation plot.

Code:

```
title: "MTCars Summary"
```

```
author: "Manasvi Sawant"
```

```
date: "2025-03-18"
```

```
output: html_document
```

```
#Intoduction
```

This report provides a summary of the mtcars dataset, including descriptive statistics and a correlation plot

```
data(mtcars)
```

```
library(ggplot2)
```


```
library(corrplot)
```

```
summary(mtcars)
library(knitr)
kable(summary(mtcars), caption = "Descriptive Statistics of mtcars Dataset")
corr_matrix<-cor(mtcars)
corrplot(corr_matrix, method = "color", type = "upper",
          rl.col="black", tl.srt = 45)
```

Output:

Table: Descriptive Statistics of mtcars Dataset

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Min.	:10.40	Min. :4.000	Min. : 71.1	Min. : 52.0	Min. :2.760	Min. :1.513	Min. :14.50	Min. :0.0000	Min. :0.0000	Min. :3.000	Min. :1.000
1st Qu.	:15.43	1st Qu.:4.000	1st Qu.:120.8	1st Qu.: 96.5	1st Qu.:3.080	1st Qu.:2.581	1st Qu.:16.89	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:3.000	1st Qu.:2.000
Median	:19.20	Median :6.000	Median :196.3	Median :123.0	Median :3.695	Median :3.325	Median :17.71	Median :0.0000	Median :0.0000	Median :4.000	Median :2.000
Mean	:20.09	Mean :6.188	Mean :230.7	Mean :146.7	Mean :3.597	Mean :3.217	Mean :17.85	Mean :0.4375	Mean :0.4062	Mean :3.688	Mean :2.812
3rd Qu.	:22.80	3rd Qu.:8.000	3rd Qu.:326.0	3rd Qu.:180.0	3rd Qu.:3.920	3rd Qu.:3.610	3rd Qu.:18.90	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:4.000	3rd Qu.:4.000
Max.	:33.90	Max. :8.000	Max. :472.0	Max. :335.0	Max. :4.930	Max. :5.424	Max. :22.90	Max. :1.0000	Max. :1.0000	Max. :5.000	Max. :8.000



30) Create a function that accepts a vector and a value. The function should return the index position of the value if found; otherwise, return "Not Found". Test it on `mtcars$mpg`.

Code:

```
find_value_index <- function(vec, value) {
  index <- which(vec == value)
  if (length(index) > 0) {
    return(index)
  } else {
    return("Not Found")
  }
}

value_to_find <- 21.0
result <- find_value_index(mtcars$mpg, value_to_find)
cat("Result for value", value_to_find, "in mtcars$mpg:", result, "\n")
```

Output:

```
> find_value_index <- function(vec, value) {
+   index <- which(vec == value)
+   if (length(index) > 0) {
+     return(index)
+   } else {
+     return("Not Found")
+   }
+ }
> value_to_find <- 21.0
> result <- find_value_index(mtcars$mpg, value_to_find)
> cat("Result for value", value_to_find, "in mtcars$mpg:", result, "\n")
Result for value 21 in mtcars$mpg: 1 2
```

Conclusion: In this experiment we learn about implementing the data structures in R on datasets. About Vectors, Lists, Dataframes, Matrices, Arrays, Factors, Tibbles data structures in R programming.

For Faculty Use

Correction Parameters	Formative Assessment [40%]	Timely completion of Practical [40%]	Attendance / Learning Attitude [20%]	
Marks Obtained				